

Mielage Prediction with Regression Analysis (Machine learning)

▼ Objective

To develop an Machine Learning Model to predict mileage using Linear Regression method

▼ Data Source

This dataset was taken from the github library which is maintained at YBI Foundation. This dataset was used in the 1983 American Statistical Association Exposition.

▼ Import Library

```
import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns
```

▼ Import Data

```
df = pd.read_csv("https://github.com/YBI-Foundation/Dataset/raw/main/MPG.csv")
```

▼ Describe Data

```
df.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	usa	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	usa	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	usa	ford torino

```
df.nunique()
```

mpg	129
cylinders	5
displacement	82
horsepower	93
weight	351
acceleration	95
model_year	13
origin	3
name	305
dtype:	int64

▼ Data Preprocessing

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    mpg         398 non-null    float64
1  cylinders    398 non-null    int64
2  displacement 398 non-null    float64
3  horsepower   392 non-null    float64
4  weight       398 non-null    int64
5  acceleration 398 non-null    float64
6  model_year   398 non-null    int64
7  origin       398 non-null    object
8  name         398 non-null    object
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB
```

```
df.describe()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
count	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	76.010050
std	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689	3.697627
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	73.000000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000	76.000000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000	79.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000



```
df.corr()
```

```
<ipython-input-10-2f6f6606aa2c>:1: FutureWarning: The default value of numeric_only in DataFrame.corr is
df.corr()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
mpg	1.000000	-0.775396	-0.804203	-0.778427	-0.831741	0.420289	0.579267
cylinders	-0.775396	1.000000	0.950721	0.842983	0.896017	-0.505419	-0.348746
displacement	-0.804203	0.950721	1.000000	0.897257	0.932824	-0.543684	-0.370164
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361
weight	-0.831741	0.896017	0.932824	0.864538	1.000000	-0.417457	-0.306564
acceleration	0.420289	-0.505419	-0.543684	-0.689196	-0.417457	1.000000	0.288137
model_year	0.579267	-0.348746	-0.370164	-0.416361	-0.306564	0.288137	1.000000



```
df = df.dropna()
```


```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 392 entries, 0 to 397
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  ---
0    mpg         392 non-null    float64
1  cylinders    392 non-null    int64
2  displacement 392 non-null    float64
3  horsepower   392 non-null    float64
4  weight       392 non-null    int64
5  acceleration 392 non-null    float64
6  model_year   392 non-null    int64
7  origin       392 non-null    object
8  name         392 non-null    object
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```


➤ Data Visualization

```
sns.pairplot(df, x_vars=['displacement', 'horsepower', 'weight', 'acceleration', 'mpg'], y_vars=['mpg'])
```

```
<seaborn.axisgrid.PairGrid at 0x7f3b3bf35180>
```



```
sns.regplot(x = 'displacement', y = 'mpg', data = df);
```



▼ Define Target Variable (y) and Feature Variables (X)

```
df.columns
```

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',  
      'acceleration', 'model_year', 'origin', 'name'],  
      dtype='object')
```

```
y = df['mpg']
```

```
y.shape
```

```
(392,)
```

```
x = df[['displacement', 'horsepower', 'weight', 'acceleration']]
```

```
x.shape
```

```
(392, 4)
```

```
x
```

	displacement	horsepower	weight	acceleration
0	307.0	130.0	3504	12.0
1	350.0	165.0	3693	11.5
2	318.0	150.0	3436	11.0
3	304.0	150.0	3433	12.0

▼ Train Test Split

Scaling Data

array([[307.0, 130.0, 3504.0, 12.0],

```
from sklearn.preprocessing import StandardScaler
```

array([[307.0, 130.0, 3504.0, 12.0],

```
ss = StandardScaler()
```

```
x = ss.fit_transform(x)
```

x

```
array([[ 1.07728956,  0.66413273,  0.62054034, -1.285258  ],
       [ 1.48873169,  1.57459447,  0.84333403, -1.46672362],
       [ 1.1825422 ,  1.18439658,  0.54038176, -1.64818924],
       ...,
       [-0.56847897, -0.53247413, -0.80463202, -1.4304305 ],
       [-0.7120053 , -0.66254009, -0.41562716,  1.11008813],
       [-0.72157372, -0.58450051, -0.30364091,  1.40043312]])
```

```
pd.DataFrame(x).describe()
```

	0	1	2	3
count	3.920000e+02	3.920000e+02	3.920000e+02	3.920000e+02
mean	-7.250436e-17	-1.812609e-16	-1.812609e-17	4.350262e-16
std	1.001278e+00	1.001278e+00	1.001278e+00	1.001278e+00
min	-1.209563e+00	-1.520975e+00	-1.608575e+00	-2.736983e+00
25%	-8.555316e-01	-7.665929e-01	-8.868535e-01	-6.410551e-01
50%	-4.153842e-01	-2.853488e-01	-2.052109e-01	-1.499869e-02
75%	7.782764e-01	5.600800e-01	7.510927e-01	5.384714e-01
max	2.493416e+00	3.265452e+00	2.549061e+00	3.360262e+00

After Standardization Mean is Zero and Standard Deviation is One

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split (x, y, train_size=0.7, random_state = 2529)
```

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

((274, 4), (118, 4), (274,), (118,))

▼ Modeling

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit (x_train, y_train)
```

▼ LinearRegression

LinearRegression()

▼ Model Evaluation

```
model.intercept_

23.485738559737584
```

```
model.coef_

array([-1.05767743, -1.68734727, -4.10787617, -0.11495177])
```

Mileage = 23.4 - 1.05Displacment + error , Mileage = 23.4 - 1.68Horsepower + error, Mileage = 23.4 - 4.10Weight + error, Mileage = 23.4 - 0.11Acceleration + error

▼ Prediction

```
y_pred = model.predict(x_test)
```

```
y_pred

array([18.51865637, 15.09305675, 14.30128789, 23.6753321 , 29.7546115 ,
       23.68796629, 26.61066644, 24.56692437, 15.06260986, 11.94312046,
       24.08050053, 27.96518468, 31.66130278, 31.01309132, 18.32428976,
       19.32795009, 28.08847536, 32.1506879 , 31.15859692, 27.15792144,
       18.82433097, 22.54580176, 26.15598115, 32.36393869, 20.74377679,
       8.78027518, 22.19699435, 18.20614294, 25.00052718, 15.26421552,
       23.13441082, 17.10542257, 9.87180062, 30.00790415, 20.41204655,
       29.11860245, 24.4305187 , 21.72601835, 10.51174626, 13.12426391,
       21.41938406, 19.96113872, 6.19146626, 17.79025345, 22.5493033 ,
       29.34765021, 13.4861847 , 25.88852083, 29.40406946, 22.41841964,
       22.07684766, 16.46575802, 24.06290693, 30.12890046, 10.11318121,
       9.85011438, 28.07543852, 23.41426617, 20.08501128, 30.68234133,
       20.92026393, 26.78370281, 22.9078744 , 14.15936872, 24.6439883 ,
       26.95515832, 15.25709393, 24.11272087, 30.80980589, 14.9770217 ,
       27.67836372, 24.2372919 , 10.92177228, 30.22858779, 30.88687365,
       27.33992044, 31.18447082, 10.8873597 , 27.63510608, 16.49231363,
       25.63229888, 29.49776285, 14.90393439, 32.78670687, 30.37325244,
       30.9262743 , 14.71702373, 27.09633246, 26.69933806, 29.06424799,
       32.45810182, 29.44846898, 31.61239999, 31.57891837, 21.46542321,
       31.76739191, 26.28605476, 28.96419915, 31.09628395, 24.80549594,
       18.76490961, 23.28043777, 23.04466919, 22.14143162, 15.95854367,
       28.62870918, 25.58809869, 11.4040908 , 25.73334842, 30.83500051,
       21.94176255, 15.34532941, 30.37399213, 28.7620624 , 29.3639931 ,
       29.10476703, 20.44662365, 28.11466839])
```

▼ Model Accuracy

```
from sklearn.metrics import mean_absolute_percentage_error, r2_score
```

```
mean_absolute_percentage_error(y_test, y_pred)

0.14713035779536746
```

```
r2_score(y_test, y_pred)

0.7031250746717691
```

▼ Explanation

mean_absolute_percentage_error = 0.14 =14% of error

therefore accuracy of our model is 100 -14 = 86%

r2_score is used to evaluate the performance of the machine learning model

here, r2_score = 0.70

which is close to 1

if r2 value is 1, it means our model is well trained for unseen data

if r2 value is 0, it means our model is badly trained for unseen data

hence our machine learning model for mileage prediction is 86% accurate

Double-click (or enter) to edit

✓

0s

completed at 6:51 PM

×